

Shaik Ibtihajulla Sha

AI Engineer | Agentic AI Developer | LLM Systems & AI Automation

+46 793229299 | shaikibtihaj07@gmail.com | Stockholm, Sweden

linkedin | github | shaik-ibtihaj

PROFESSIONAL SUMMARY

Applied AI Engineer with hands-on experience designing and deploying agentic AI systems, LLM-powered workflows, and production-grade automation pipelines. Proficient in Python, REST APIs, and modern AI engineering frameworks — including LLM orchestration, tool calling, RAG, multi-agent systems, and prompt engineering. Demonstrated ability to bridge technical depth and business impact: translating client requirements into reliable AI solutions that reduce manual effort and create measurable operational value. Well-suited to consulting environments where agentic AI architecture, stakeholder communication, and end-to-end ownership are equally important.

EDUCATION

Blekinge Institute of Technology (BTH), Sweden

2025 – 2026

Exchange Program in Computer Science

Relevant Coursework: Machine Learning, Applied Statistics, Data Analytics, Statistical Modeling, Databases & SQL, Software Engineering, Python Programming

EXPERIENCE

Data Analyst Intern

AxioGreen — Energy & Sustainability Tech

Python · SQL · REST APIs · Data Pipelines · Workflow Automation · Stakeholder Collaboration 3–6 months

- Automated data collection workflows using Python and REST API integrations, eliminating manual data handling and reducing operational overhead across energy monitoring processes.
- Built and maintained end-to-end data pipelines ingesting energy consumption metrics from backend systems into reporting layers, ensuring reliable, low-latency data availability for business stakeholders.
- Developed analytical dashboards visualizing key energy KPIs, enabling non-technical operations leaders to monitor performance trends and act on data-driven insights independently.
- Collaborated cross-functionally with engineering and product teams to translate business requirements into technical solutions, communicating findings and tradeoffs clearly to non-technical audiences.
- Took end-to-end ownership of multiple data automation initiatives — from requirements discovery through deployment — demonstrating a consulting mindset and proactive stakeholder engagement.

PROJECTS

Agentic AI Workflow Assistant

Python · LLM Orchestration · Tool Calling · OpenAI API · REST APIs · Multi-Step

Reasoning · AI Automation

- Architected a production-oriented agentic AI system leveraging LLM orchestration and function calling to autonomously execute multi-step business workflows, reducing manual decision overhead by automating repetitive operational tasks.
- Implemented tool-calling integrations with external REST APIs, enabling the agent to retrieve real-time data, trigger downstream actions, and produce context-aware outputs without human intervention at each step.
- Designed a modular multi-agent architecture with dedicated planning, execution, and validation stages — supporting agent collaboration, structured reasoning, and reliable task decomposition across complex workflows.
- Built reliability mechanisms including retry logic, fallback routing, and structured error handling to ensure consistent, production-grade performance under real-world failure conditions.
- Developed a lightweight evaluation framework to measure agent task completion accuracy and output quality, enabling iterative prompt engineering and system refinement.

Enterprise RAG Knowledge Assistant

Python · LangChain · OpenAI API · Vector Database · RAG · Embeddings ·

Docker

- Built an enterprise-grade knowledge assistant using Retrieval-Augmented Generation (RAG), processing internal documentation corpora into semantically indexed vector database stores for high-precision information retrieval.
- Implemented embedding-based semantic search to surface contextually relevant document chunks, enabling the assistant to generate context-aware, citation-grounded answers — reducing information retrieval time significantly.
- Designed a document ingestion and chunking pipeline with metadata filtering, supporting large-scale knowledge bases and ensuring retrieval quality across diverse document types.
- Integrated evaluation metrics (retrieval precision, answer faithfulness) to benchmark system performance and guide iterative improvements to chunking strategy and prompt design.
- Containerized the full system with Docker for production-ready deployment, with modular architecture supporting integration into existing enterprise workflows via REST API.

Multi-Agent Business Process Automation Platform

Python · OpenAI Agents SDK · MCP · Tool Calling · Docker · Monitoring

- Architected a multi-agent automation platform using the OpenAI Agents SDK, with specialized agents responsible for planning, execution, and validation — enabling autonomous end-to-end business workflow execution.
- Integrated Model Context Protocol (MCP) to connect agents with external tools and enterprise data sources, enabling dynamic context injection and real-time data retrieval without hardcoded tool configurations.
- Designed human-in-the-loop validation checkpoints for high-stakes decision nodes, balancing autonomous execution with appropriate oversight in business process automation scenarios.
- Implemented agent observability and monitoring pipelines tracking task completion rates, latency, and error distribution — supporting reliability analysis and production incident response.
- Delivered measurable automation outcomes: multi-agent collaboration reduced workflow execution time from hours to seconds on targeted business processes.

Probabilistic User Behavior Modeling — Bachelor’s Thesis

Python · Markov Models · Statistical Analysis · Synthetic Data Generation · Behavioral Analytics

- Designed first- and second-order Markov chain models to probabilistically model sequential user interaction patterns, uncovering latent behavioral structure applicable to product analytics and AI personalization systems.
- Generated synthetic session datasets for privacy-safe experimentation and conducted rigorous comparative model evaluation, documenting methodology and findings in a full academic thesis.

Customer Churn Prediction System

Python · Scikit-Learn · Pandas · SQL · Feature Engineering · Model Evaluation

- Engineered an end-to-end ML pipeline — data cleaning, feature engineering, model training, and evaluation — achieving strong AUC-ROC improvement over baseline on high-risk churn classification targets.
- Translated model outputs into stakeholder-facing dashboards and narrative summaries, linking predictive insights directly to actionable retention strategy recommendations.

TECHNICAL SKILLS

Programming	Python, SQL, JavaScript
AI Engineering	LLMs, Agentic AI, OpenAI API, Anthropic API (Claude), Prompt Engineering, Tool Calling, Function Calling, RAG, Vector Databases, Embeddings, Multi-Agent Systems, MCP, AI Evaluation
Machine Learning	Scikit-Learn, Pandas, NumPy, Model Development, Feature Engineering, Statistical Modeling
Backend & APIs	Flask, FastAPI, REST APIs, Microservices
Cloud & DevOps	Azure, Docker, Git, Linux, CI/CD
Frameworks & Tools	LangChain, OpenAI Agents SDK, Firebase, Node.js, React

CORE COMPETENCIES

Agentic AI Development | LLM Orchestration | AI Automation | Multi-Agent Systems | Production Deployment | Reliability Engineering | RAG & Knowledge Systems | Prompt Engineering | Client Communication | Stakeholder Management | Business Analysis | Consulting Mindset | Cross-Functional Collaboration

LANGUAGES

English — Fluent | Swedish — Intermediate (actively improving)